

# Time and Ordering of Events

Dominic Duggan  
Stevens Institute of Technology  
Partly based on material by  
Sape Mullender and Ken Birman

1

1

## What time is it?

- Agree that update A occurred before update B
- Offer a “lease” on a resource that expires at time 10:10.0150
- Guarantee that a time critical event will reach all interested parties within 100ms

2

2

## What does time “mean”?

- Time on a global clock?
  - E.g. with GPS receiver
- Machine’s local clock
  - But was it set accurately?
  - And could it drift, e.g. run fast or slow?
  - What about faults, like stuck bits?
- Or try to agree on time

3

3

## LOGICAL TIME

4

4


## Lamport: Logical Time

- Time lets a system ask “Which came first: event A or event B?”
- Time is a means of labeling events so that...
  - If A happened before B,  $TIME(A) < TIME(B)$
  - If  $TIME(A) < TIME(B)$ , A happened before B

5

5


## Drawing time-line pictures:



6

6

## Drawing time-line pictures:




- A, B, C and D are “events”.
  - So are `snd(m)` and `rcv(m)` and `deliv(m)`
- What ordering claims are meaningful?

7

7

## Drawing time-line pictures:




- A happens before B, and C before D
  - *Local ordering* at a single process
  - Write  $A \rightarrow^p B$  and  $C \rightarrow^q D$

8

8

## Drawing time-line pictures:




- $\text{snd}(m)$  also happens before  $\text{rcv}(m)$ 
  - *Distributed ordering* introduced by a message
  - Write  $\text{snd}(m) \rightarrow \text{rcv}(m)$

9

9

## Drawing time-line pictures:




- $A$  happens before  $D$ 
  - *Transitivity*:  $A$  happens before  $\text{snd}(m)$ , which happens before  $\text{rcv}(m)$ , which happens before  $D$

10

10

## Drawing time-line pictures:



- B and D are concurrent
  - Looks like B happens first, but D has no way to know. No information flowed...

11

11

## “Happens before” relation

- We'll say that “*A happens before B*”, written  $A \rightarrow B$ , if
  - $A \rightarrow^p B$  according to the local ordering, or
  - A is  $\text{snd}(m)$  and B is  $\text{rcv}(m)$  and  $A \rightarrow B$ , or
  - A and B are related under the transitive closure of rules (1) and (2)

12

12

## LOGICAL CLOCKS

13

13

## Logical clocks

- First version: uses just a single integer
  - Designed for big (64-bit or more) counters
  - Each process  $p$  maintains  $LT_p$ , a local counter
  - A message  $m$  will carry timestamp  $TS(m)$

14

14

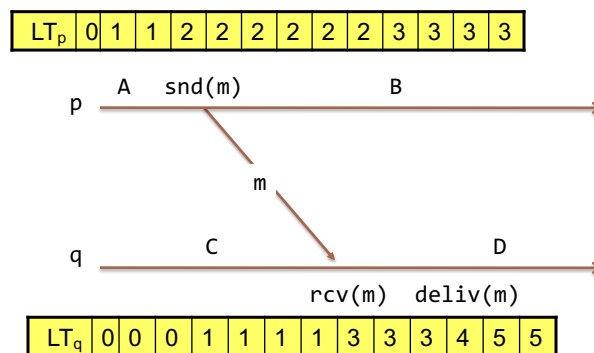
## Rules for managing logical clocks

- When an event happens at a process  $p$  it increments  $LT_p$ 
  - Any event that matters to  $p$
  - Normally, also  $snd$  and  $rcv$  events (since we want receive to occur “after” the matching send)
- When  $p$  sends  $m$ , set
  - $TS(m) = LT_p$
- When  $q$  receives  $m$ , set
  - $LT_q = \max(LT_q, TS(m)) + 1$

15

15

## Time-line with LT annotations




- $LT(A) = 1$ ,  $LT(snd(m)) = 2$ ,  $TS(m) = 2$
- $LT(rcv(m)) = \max(1, 2) + 1 = 3$ , etc...


16

16








17




18







21





22




23




24





27



28

## **DISTRIBUTED MUTUAL EXCLUSION**

29

29

## Distributed Mutual Exclusion

- Idea: purely distributed protocol for mutually exclusive access to a resource
  - No central coordinator

30

30

## Distributed Mutual Exclusion

- Idea: purely distributed protocol for mutually exclusive access to a resource
  - No central coordinator
- Requests are ordered using logical time
  - Use (ts, pid) with pid to break ties
  - $(m, p) < (n, q)$  if  $m < n$  or  $(m = n \text{ and } p < q)$

31

31

## Distributed Mutual Exclusion

- Idea: purely distributed protocol for mutually exclusive access to a resource
  - No central coordinator
- Requests are ordered using logical time
  - Use (ts, pid) with pid to break ties
  - $(m, p) < (n, q)$  if  $m < n$  or  $(m = n \text{ and } p < q)$
- Data structures
  - Logical time  $LT_p$
  - Request queue, ordered by request timestamp
  - $LT[i]$ , timestamp of last message received from process  $p_i$

32

32



## Request a Resource

- Process  $p_i$ 
  - Increments its logical clock
  - Adds request  $m$  ( $TS(m)=LT_i$ ) to its request queue
  - Broadcasts request to every other process

33

33

## Request a Resource

- Process  $p_i$ 
  - Increments its logical clock
  - Adds request  $m$  ( $TS(m)=LT_i$ ) to its request queue
  - Broadcasts request to every other process
- Process  $p_j$  ( $j \neq i$ )
  - Acknowledges receipt of request ( $TS(ack)=LT_j$ )

34


34

## Request a Resource

- Process  $p_i$ 
  - Increments its logical clock
  - Adds request  $m$  ( $TS(m)=LT_i$ ) to its request queue
  - Broadcasts request to every other process
- Process  $p_j$  ( $j \neq i$ )
  - Acknowledges receipt of request ( $TS(ack)=LT_j$ )
- Process  $p_i$  has access when:
  - Its request is in the front of its request queue
  - $LT[i] \geq TS(m)$  for all  $i=1,...,n$

35

35



36

36

